


# Longest-Context End-to-End Latency Summary


Baseline FlashAttention CustomSA

## GPU FP16 @ 32k




Latency (ms)

## AWS FP16 @ 32k




Latency (ms)

## TPU FP16 @ 32k




Latency (ms)

## GPU INT8 @ 32k (No Conversion)




Latency (ms)

## AWS INT8 @ 32k (No Conversion)



Latency (ms)

## TPU INT8 @ 32k (No Conversion)



Latency (ms)